

---

# NoSQL Databases



---

# What are SQL databases?

Relational Databases

... Rows & Columns



---

# What are SQL databases?

- Structured Data Management
- Tables with rows (records) and columns (fields)

Col1	Col2	Col3	Col4
A	—	—	—
B	—	—	—

- SQL (Structured Query Language) for data manipulation



---

# Moving Beyond Relational Databases

- Relational databases (RDBMS) have ruled for decades.
- Big Data brings new challenges:
  - **Volume, Velocity, Variety**
- Rigid schemas can become a bottleneck.

---

# What is NoSQL?



---

# What is NoSQL?

NoSQL stands for "Not Only SQL"

(but sometimes means "Non-SQL")



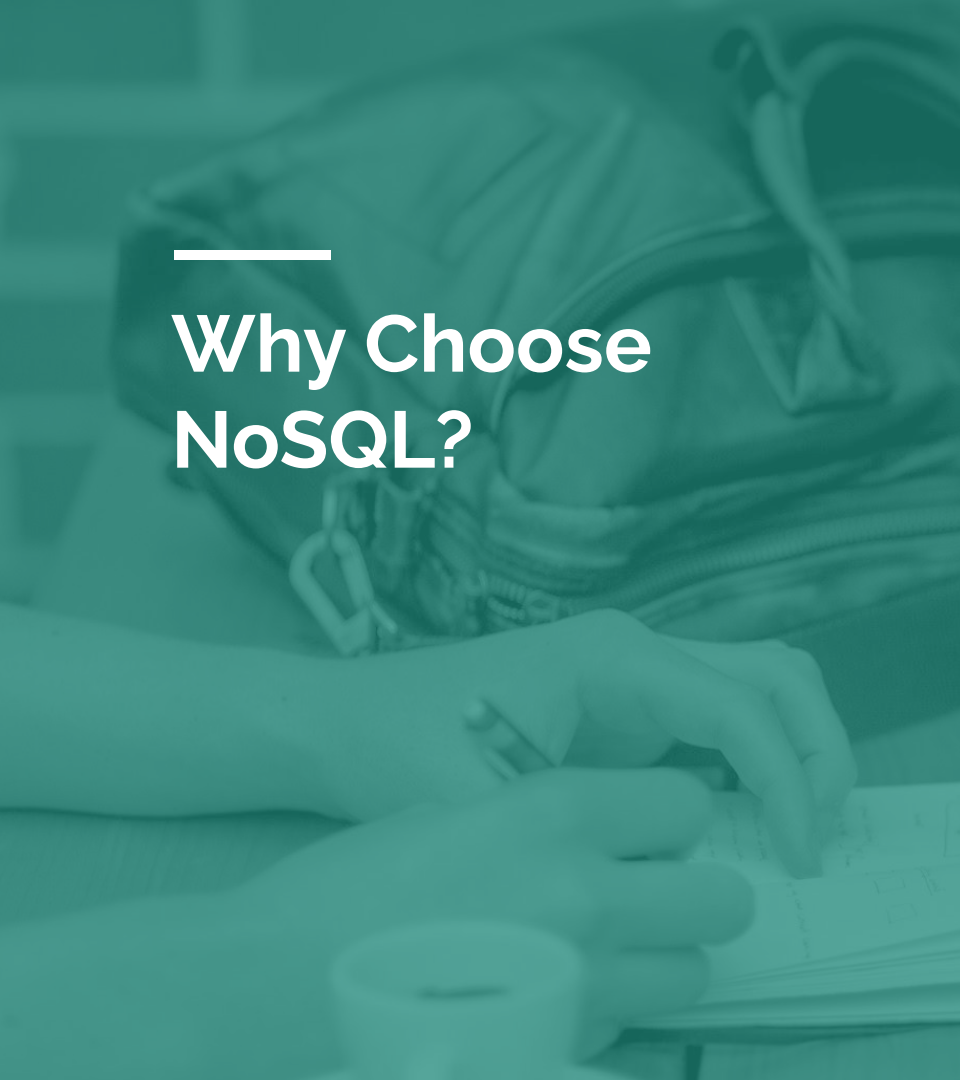
---

# What is NoSQL?

NoSQL stands for "Not Only SQL"

(but sometimes means "Non-SQL")

An umbrella term for databases that don't use SQL as their primary query language



---

# Why Choose NoSQL?

## Why Choose NoSQL?

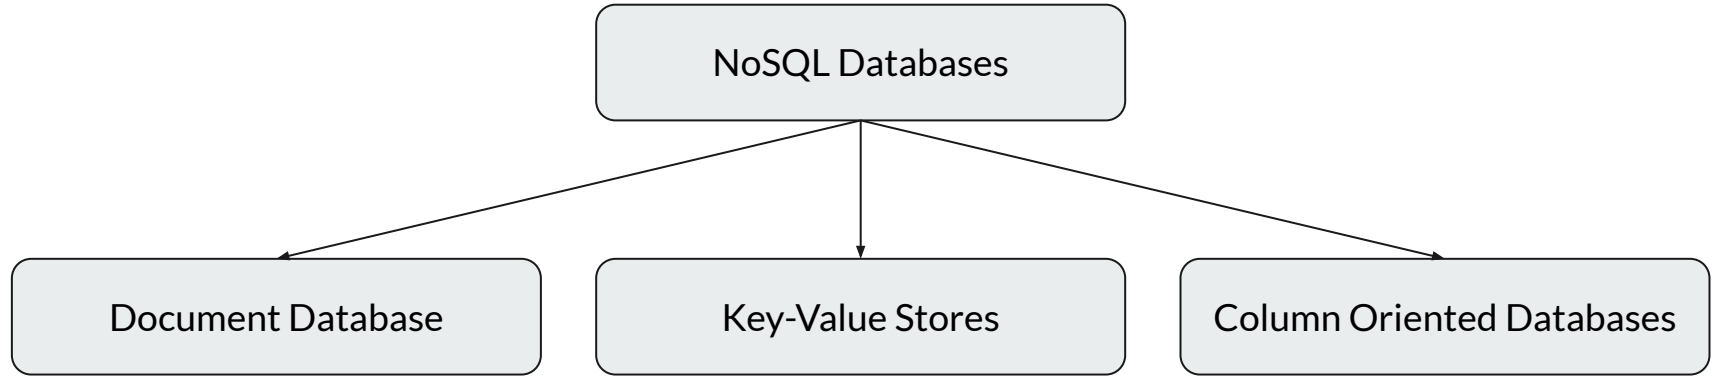
- Handling massive, semi-structured or unstructured data
- Flexibility for evolving data models.



---

# Types of NoSQL Databases

# Types of NoSQL Databases



---

# Document Databases

Document Databases: Flexibility  
for Complex Data





# Document Databases

## Document Databases: Flexibility for Complex Data

- Store data in documents (flexible structures like JSON)
- Documents can have nested structures (like profiles with addresses)

```
{
  "_id": "5cf0029caff5056591b0ce7d",
  "firstname": "Jane",
  "lastname": "Wu",
  "address": {
    "street": "1 Circle Rd",
    "city": "Los Angeles",
    "state": "CA",
    "zip": "90404"
  },
  "hobbies": ["surfing", "coding"]
}
```



# Document Databases

## Document Databases: Flexibility for Complex Data

- Store data in documents (flexible structures like JSON)
- Documents can have nested structures (like profiles with addresses)
- Ideal for storing semi-structured or unstructured data



---

# Key-Value Stores

## Key-Value Stores: Fast Lookups for Simple Data

- Simplest NoSQL model: key-value pairs.
  - Keys are unique identifiers (like user IDs, product codes)
  - Values can be various data types (strings, numbers, JSON)
- Ideal for fast retrievals based on the key

---

# Key-Value Stores

## Redis

- In-memory database known for lightning-fast performance.
- Used by Craigslist, **Instagram**, StackOverflow, flickr, ...

Instagram needed to keep around a **mapping** of about 300 million photos back to the **user ID** that created them.



---

# Column-Oriented Databases

## Column-Oriented Databases: Built for Speed and Analytics

- Store data by columns instead of rows (unlike relational databases)
- Optimized for reading specific columns – faster analytics
- Handle schema changes more efficiently (add new columns easily)



# Information storage in Column-Oriented Storage

Product ID	Product Name	Price	Quantity
------------	--------------	-------	----------

1	Widget Deluxe	24.99	50
---	---------------	-------	----

2	Super Gadget	15.50	20
---	--------------	-------	----

3	Shiny Thing	9.99	100
---	-------------	------	-----

But then how do I access all  
info for **product\_id=1**?

## Traditional Relational (Row-Oriented) Storage

- **Row 1:** 1, Widget Deluxe, 24.99, 50
- **Row 2:** 2, Super Gadget, 15.50, 20
- **Row 3:** 3, Shiny Thing, 9.99, 100

## Column-Oriented Storage (like Bigtable)

- **Product ID Column:** 1, 2, 3
- **Product Name Column:** Widget Deluxe, Super Gadget, Shiny Thing
- **Price Column:** 24.99, 15.50, 9.99
- **Quantity Column:** 50, 20, 100

# Information storage in Column-Oriented Storage

Product ID	Product Name	Price	Quantity
1	Widget Deluxe	24.99	50
2	Super Gadget	15.50	20
3	Shiny Thing	9.99	100

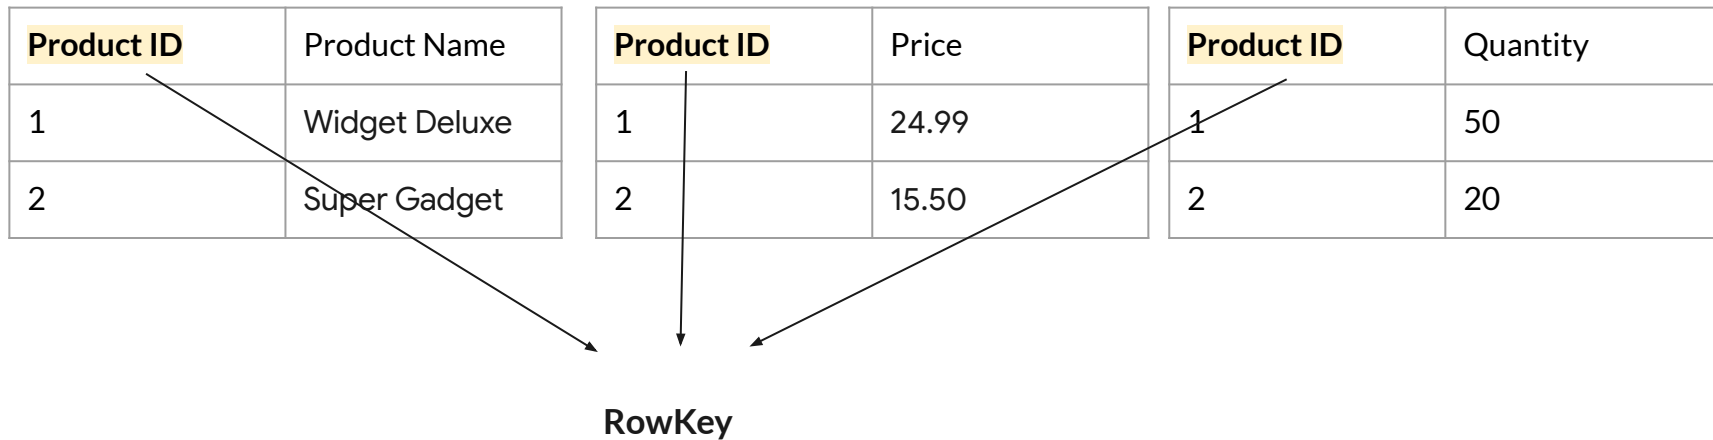
Replication

Product ID	Product Name
1	Widget Deluxe
2	Super Gadget

Product ID	Price
1	24.99
2	15.50

Product ID	Quantity
1	50
2	20

# Information storage in Column-Oriented Storage



Row keys offer a mechanism to uniquely identify and logically group data.

---

# Column-Oriented Databases

## Google's Bigtable



### Google's Bigtable: A NoSQL Pioneer for Massive Data

- Column-oriented, distributed NoSQL database developed by Google.
- Designed to handle petabytes of data.
- Powers many Google services (Search, Maps, Gmail)

---

# Column-Oriented Databases

## Google's Bigtable



### Google's Bigtable: A NoSQL Pioneer for Massive Data

- Column-oriented, distributed NoSQL database developed by Google.

Bigtable's success inspired creation of HBase!

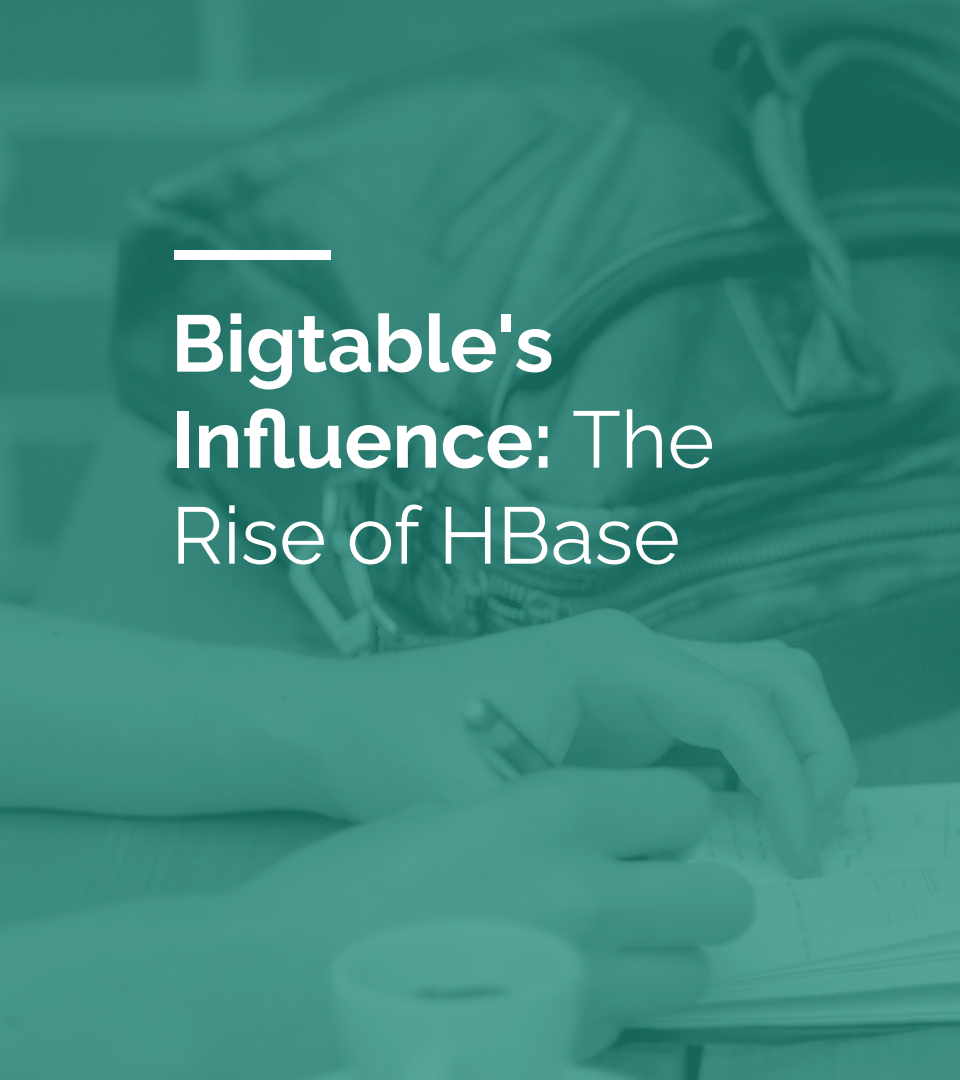
---

# The Rise of HBase

---

# Bigtable's Influence: The Rise of HBase





---

# Bigtable's Influence: The Rise of HBase

## Google's Bigtable: A NoSQL Pioneer for Massive Data

- Google published a paper (2006) describing Bigtable's design and concepts.
- This sparked interest in column-oriented NoSQL solutions outside of Google.
- HBase emerged as an open-source implementation heavily inspired by Bigtable



---

# **Bigtable and HBase**

## A Comparative Look



---

# Bigtable and HBase

## A Comparative Look

### Google's Bigtable: A NoSQL Pioneer for Massive Data

- Similarities:
  - Column-oriented data model
  - Sparse table structure
  - Built for scalability and high performance.
- Differences:
  - Bigtable runs on Google's infrastructure; HBase runs on Hadoop.
  - HBase has stronger consistency guarantees.
  - Bigtable's proprietary API vs. HBase's more open interfaces.



APACHE  
**HBASE**

# Welcome to the Zoo!





---

# HBase: Scalable, Distributed Storage on Hadoop


## HBase: Scalable, Distributed Storage on Hadoop

- Open-source, column-oriented NoSQL database.
- Part of the Apache Hadoop ecosystem, built on top of HDFS.
- Designed for large-scale, structured and semi-structured data.
- Real-time read/write access.

---

# HBase Architecture


# HBase Architecture



Apache HBase Architecture

# HMaster

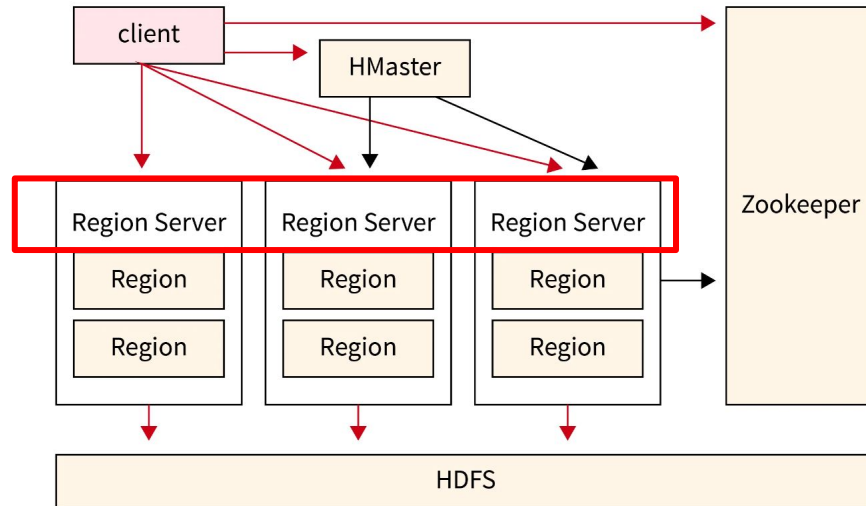
- Manages the cluster's RegionServers.
- Assigns Regions (chunks of data) to RegionServers.
- Responsible for load balancing across the cluster.
- Monitors RegionServer health and handles failures.
- Handles metadata operations (table creation, etc.)





# RegionServers

- Serve data for read and write requests.
- Manage Regions, which are segments of an HBase table.
- Split Regions that grow too large.
- Handle data compactions for storage efficiency.

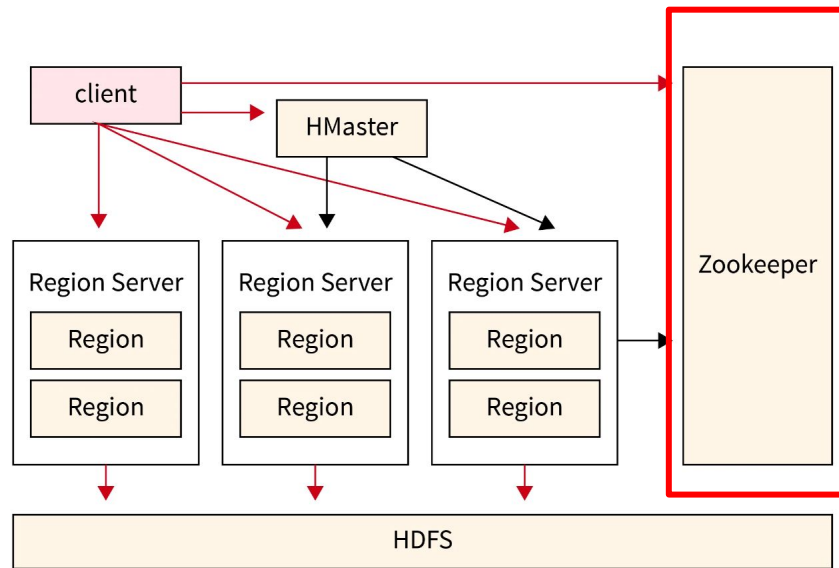


Apache HBase Architecture

# Zookeeper

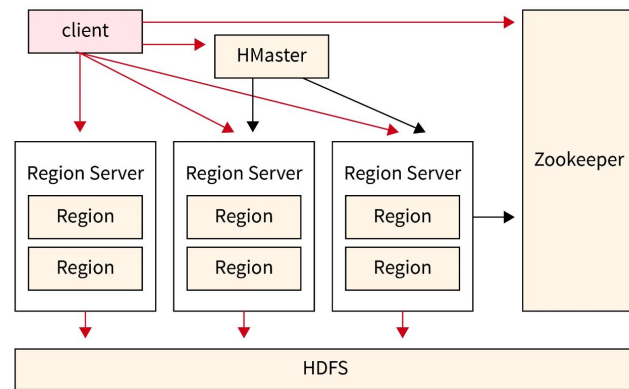
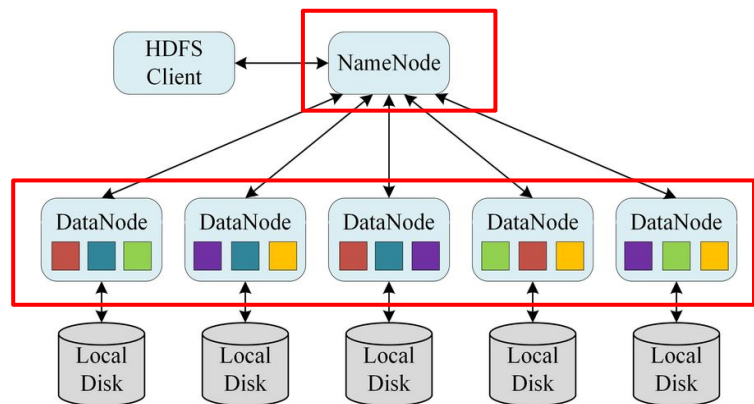
- A coordination service essential for distributed systems.

Have we seen a similar  
Master-Slave storage setup  
before?



Apache HBase Architecture

# Recall ...



Apache HBase Architecture

## HBase vs HDFS?



- **NameNode:** A **librarian** keeping track of which books are on which shelves in the library.
- **HMaster:** The manager of a specific section of the library (e.g., the fiction section), deciding where to put new books, rearranging shelves, and handling requests from readers within that section.

---

# HBase: Scalable, Distributed Storage on Hadoop

## HBase vs HDFS?

- Layer of Operation:
  - **HMaster** operates at the level of HBase, managing its specific data structures and requests.
  - **NameNode** operates at the level of HDFS, responsible for the file system itself.
- NameNode primarily handles metadata (data about data).
- HMaster manages metadata about HBase tables
- HBase depends on HDFS for underlying storage.

---

# HBase Data Model

# Sample Database

Row Key	info:Browser	info:DeviceType	info:Location	actions:PageVisited	actions:SearchTerm	actions:TimeOnPage (sec)
UserID#1234_202311271230 (T1)	Chrome	Desktop	New York, USA	/home	product reviews	125
UserID#1234_202311271235 (T2)	Chrome	Desktop	New York, USA	/products/gadgets		68
UserID#5678_202311262310 (T1)	Safari	Mobile	London, UK	/about	company history	42
UserID#9012_202311270945 (T1)	Firefox	Tablet	Berlin, DE	/home		27
UserID#9012_202311270950 (T2)	Firefox	Tablet	Berlin, DE	/products/software	antivirus deals	85

# Sample Database

Row Key	info:Browser	info:DeviceType	info:Location	actions:PageVisited	actions:SearchTerm	actions:TimeOnPage (sec)
UserID#1234_202311271230 (T1)	Chrome	Desktop	New York, USA	/home	product reviews	125
UserID#1234_202311271235 (T2)	Chrome	Desktop	New York, USA	/products/gadgets		68
UserID#5678_202311262310 (T1)	Safari	Mobile	London, UK	/about	company history	42
UserID#9012_202311270945 (T1)	Firefox	Tablet	Berlin, DE	/home		27
UserID#9012_202311270950 (T2)	Firefox	Tablet	Berlin, DE	/products/software	antivirus deals	85

RowKey: Unique identifier  
for each row.



# Sample Database

Row Key	info:Browser	info:DeviceType	info:Location	actions:PageVisited	actions:SearchTerm	actions:TimeOnPage (sec)
UserID#1234_202311271230 (T1)	Chrome	Desktop	New York, USA	/home	product reviews	125
UserID#1234_202311271235 (T2)	Chrome	Desktop	New York, USA	/products/gadgets		68
UserID#5678_202311262310 (T1)	Safari	Mobile	London, UK	/about	company history	42
UserID#9012_202311270945 (T1)	Firefox	Tablet	Berlin, DE	/home		27
UserID#9012_202311270950 (T2)	Firefox	Tablet	Berlin, DE	/products/software	antivirus deals	85

Column Families: Groups of related columns.

# Sample Database

Row Key	info:Browser	info:DeviceType	info:Location	actions:PageVisited	actions:SearchTerm	actions:TimeOnPage (sec)
UserID#1234_202311271230 (T1)	Chrome	Desktop	New York, USA	/home	product reviews	125
UserID#1234_202311271235 (T2)	Chrome	Desktop	New York, USA	/products/gadgets		68
UserID#5678_202311262310 (T1)	Safari	Mobile	London, UK	/about	company history	42
UserID#9012_202311270945 (T1)	Firefox	Tablet	Berlin, DE	/home		27
UserID#9012_202311270950 (T2)	Firefox	Tablet	Berlin, DE	/products/software	antivirus deals	85

Columns

# Sample Database

Row Key		info:Browser	info:DeviceType	info:Location	actions:PageVisited	actions:SearchTerm	actions:TimeOnPage (sec)
UserID#1234_	202311271230 (T1)	Chrome	Desktop	New York, USA	/home	product reviews	125
UserID#1234_	202311271235 (T2)	Chrome	Desktop	New York, USA	/products/gadgets		68
UserID#5678_	202311262310 (T1)	Safari	Mobile	London, UK	/about	company history	42
UserID#9012_	202311270945 (T1)	Firefox	Tablet	Berlin, DE	/home		27
UserID#9012_	202311270950 (T2)	Firefox	Tablet	Berlin, DE	/products/software	antivirus deals	85

Timestamp:  
To keep track  
of changes

# Sample Database

Row Key	info:Browser	info:Device	actions:TimeOnPage (sec)
UserID#1234_202311271230 (T1)	Chrome	Desktop	125
UserID#1234_202311271235 (T2)	Chrome	Desktop	68
UserID#5678_202311262310 (T1)	Safari	Mobile	42
UserID#9012_202311270945 (T1)	Firefox	Tablet	27
UserID#9012_202311270950 (T2)	Firefox	Tablet	85

RowKey = UserID + Timestamp

RowKey: Unique identifier for each row.



---

# RowKey Design

## The Importance of Row Key Design

- **Distribution:** Row keys determine how data is spread across the cluster.
- **Avoid hotspots:** Poor row key design can lead to concentrated load on certain machines.
- **Consider access patterns:** How will you query your data?



# RowKey Design

## The Importance of Row Key Design

- HBase stores data **sorted by Row Key**. This enables:
  - Fast Lookups: If you know the Row Key, HBase can quickly navigate to the correct data location.
  - Range Scans: Efficiently retrieves sets of rows within a key range.

# RowKey Design

## Row Key

UserID#1234_202311271230 (T1)
UserID#1234_202311271235 (T2)
UserID#5678_202311262310 (T1)
UserID#9012_202311270945 (T1)
UserID#9012_202311270950 (T2)

RowKey: Unique identifier  
for each row.

## The Importance of Row Key Design

- **Consider access patterns:** How will you query your data?

Think about how you'll be querying your data. Will you often fetch by:

- UserID?
- Date or Time Range?
- Location?
- Some other attribute? Your Row Key should make those common queries efficient.



---

# RowKey Design

## The Importance of Row Key Design

- **Distribution:** Row keys determine how data is spread across the cluster.
- **Avoid hotspots:** Poor row key design can lead to **concentrated load** on certain machines.
- **Consider access patterns:** How will you query your data?



# RowKey Design

## Row Key

UserID#1234_202311271230 (T1)
UserID#1234_202311271235 (T2)
UserID#5678_202311262310 (T1)
UserID#9012_202311270945 (T1)
UserID#9012_202311270950 (T2)

RowKey: Unique identifier  
for each row.

## The Importance of Row Key Design

- Avoid **hotspots**: Poor row key design can lead to **concentrated load** on certain machines.

Why add timestamp to the rowkey?

If our Row Key was simply the UserID?

# RowKey Design

## Row Key

UserID#1234_202311271230 (T1)
UserID#1234_202311271235 (T2)
UserID#5678_202311262310 (T1)
UserID#9012_202311270945 (T1)
UserID#9012_202311270950 (T2)

RowKey: Unique identifier  
for each row.

## The Importance of Row Key Design

### If our Row Key was simply the UserID?

- **Hotspots:** Popular users cause load imbalances.
- **Sequential Nature:** New user data concentrates on a few machines.
- **Limited Querying:** Difficult to analyze data based on time ranges.

# RowKey Design

## Row Key

UserID#1234_202311271230 (T1)
UserID#1234_202311271235 (T2)
UserID#5678_202311262310 (T1)
UserID#9012_202311270945 (T1)
UserID#9012_202311270950 (T2)

RowKey: Unique identifier  
for each row.

## The Importance of Row Key Design

How does adding **timestamp** to UserID help?

- **Distributes Recent Data:** Actions from the same user are spread out.
- **Enables Time-Based Queries:** Fetch data within time ranges and analyze temporal trends.

---

# Common HBase Commands

# Creating and Manipulating Tables



- **create** 'table\_name', 'col\_family1', 'col\_family2' ... (Creates a new table)
- **list** (Shows all tables)
- **describe** 'table\_name' (Displays table's schema)
- **disable** 'table\_name' (Disables a table)
- **drop** 'table\_name' (Deletes a table)

# Working with Data



- `put 'table_name', 'row_key', 'col_family:column', 'value'` (Inserts data)
- `get 'table_name', 'row_key'` (Retrieves data by row key)
- `scan 'table_name'` (Scans the contents of a table – use with caution on large tables)
- `delete 'table_name', 'row_key'` – (Deletes a row)